

D1 is won where *session one ends*.

Every new user walks away with a finished lesson and no promise of tomorrow. We add one hand-off at the end of the first session — a concrete “tomorrow’s five minutes” plan, a reminder at the hour they choose, and a day-2 card waiting for them. That’s the gap between 16% and 22%.

16%

D1 TODAY


+6pp

GAP TO BRIDGE

22%

INDUSTRY BENCHMARK

16% • today



6 percentage points but ~37% more of the people who tried us. That is not a content problem — it's a *returns-problem*: nothing in the product gives a user a reason to come back within 24 hours.

The day-2 return is decided at the end of day 1.

D1 is the share of new users who return the *calendar day* after their first session. The decision to come back is made while session 1 is still open — then the app goes silent for 24 hours. The few products that hit 22% don't rely on goodwill; they plant the reason to return before the user leaves.

1

Discover

A short lesson, headline, or interview prep promise.



The hand-off

Where 16% vs 22% is decided — currently nothing.

4

Day 2

Rarely comes. Today: only 16% of users.

2

Session 1

A 5-minute lesson, some news, a browse.

3

24h silence

Attention fades; the calendar day passes.

Four testable hypotheses — all live at the end of session 1.

H1 • NO HAND-OFF

Done means done

The lesson ends and so does the session. No screen ever says “here’s what’s next.”

H2 • NO STAKE

No commitment

The user hasn’t set a goal or a time slot, so nothing anchors the return ahead of the next session.

H3 • NO NUDGE

Quiet 24 hours

There’s no prompt when the user has a moment — the day passes in silence.

H4 • NO PROGRESS

Nothing carried

Finished-and-forgotten: no saved state makes day 2 worth starting (a review deck, a streak, a plan).

Each hypothesis maps to a cheap, independently-measurable intervention. We pick the two that *compound*: the hand-off and the commitment behind it (H1 + H2), delivered as one redesigned final screen — because building a hook without a reason-to-return is just decoration, and a reason without a hook never gets seen.

“Tomorrow’s five minutes.” A hand-off, not a gimmick.

We change what session 1 ends with. Instead of “thanks for learning,” the user gets a plan card: what they learned today, what tomorrow’s 5-minute lesson is, and a reminder at the hour they choose. One loop, owned end to end.

SESSION 1 ENDS

“You learned 3 terms.”

Progress is made visible and concrete — not a score, the actual knowledge.



THE PLAN CARD

“Tomorrow: Prompting basics.”

Tomorrow named + “remind me at 7:45am.” The slot is the commitment.



THE REMINDER

“Your 5 min are ready.”

Email-first for speed and reliability; push slides in where permissions exist.



DAY 2 OPENS

Back, with a plan.

The card opens the promised 5-minute lesson; the chain can repeat tomorrow.

One screen does the work.

Two small additions — a goal question at sign-up, and a redesigned final screen — turn “done” into “see you tomorrow.” Companion interactive prototype included.

Unrot

Lesson 1 of 5

Today you learned — **context windows**, **temperature**, and **token limits**. Saved to your review deck.

TOMORROW'S PLAN

Prompting basics · 5 min

Today's terms become tomorrow's prompt practice. 6 new terms, one quiz.

Remind me at 7:45 am · weekdays

Set reminder & done

Skip — I'll remember

Unrot

7:45 am · reminder

“Your 5 minutes are ready.” Today's lesson: Prompting basics.

You're back

STREAK • DAY 2

Prompting basics

Three examples of weak versus strong prompts — then a 10-second rewrite drill.

Pick up where you left off — 4 min

Start lesson

The card lives at the top of home until done — one tap from anywhere on day 2.

Why this lever wins the four-week budget.

LEVER	IMPACT ON D1	EFFORT	CONFIDENCE
End-of-session hand-off + reminder Plan card, chosen slot, email/push nudge, day-2 landing.	High	Low-Med	High
Streak + daily goal Duolingo-style chain wired to a self-set “improve at X” goal.	Med-High	Medium	Medium
Personalization / feed work Ranking lessons and news to the user’s interests.	Med	High	Low
Surface 1 question added to onboarding, 1 screen redesigned, 1 scheduler, 1 day-2 landing. Fits 2 engineers.			
Plays forward The slot and goal become the backbone for streaks and personalization later — nothing is wasted.			
Measurable D1 of the variant cohort is a single, pre-registered number. Clean read in a 50/50 split.			

The streak is deliberately *not* in the four-week spec: it multiplies the hand-off, and it lands with more confidence once we know the hand-off works. Scope discipline is the deliverable.

Four weeks, one loop, cadence with room.

Week 1 INSTRUMENT	Analyt Design Add events: <code>session1_complete</code>, <code>reminder_accept</code>, <code>day2_start</code>. Pull the session-1 funnel to confirm where users drop. Designer locks hi-fi flows for the hand-off and day-2 landing; copy written and tested on 5 users.
Week 2 BUILD	Eng Design Build the plan card + slot capture behind a feature flag. Ship to an internal beta on the final day. Funnel check: does the extra screen dent session-1 completion? (Guardrail metric live.)
Week 3 DELIVER	Eng Scheduler + email-first reminder; day-2 landing page. QA the calendar-day definition and timezones — D1 is per-user local day, so this is where silent bugs hide. Push delivery wired but gated off.
Week 4 LAUNCH	Analyt Eng 50/50 randomized rollout on all new users, code frozen. First D1 readout lands the day after the early cohort's day-2 — a directional read within the window, full power as cohorts accumulate.


Roles: the designer front-loads (weeks 1–2), both engineers carry weeks 2–4. Buffer kept for reminder deliverability — the top failure risk — with email-first as the fallback that needs no permission capture at all.

One number decides it; four protect it.

METRIC PLAN

PRIMARY	D1 retention — variant vs. control Difference in return rates for new-user cohorts, 50/50 split. Pre-registered; reported with a confidence interval, no peeking.
GUARDRAIL	Session-1 completion & logged-in session length Must not regress — the hand-off can't cost us the first session to win the second.
SECONDARY	Reminder accept → click-through → day-2 start Diagnoses where in the loop we leak: is it the card, the nudge, or the content?
COHORT	Split by entry point: lesson / news / interview prep One-shot reference users likely return differently from daily learners — informs the next quarter.

WHAT A WIN LOOKS LIKE



Ship the hand-off. Learn from the return.

Six points separate us from the benchmark, and none of them require new content — they require treating the end of session 1 as a decision point instead of a goodbye.

The solution

Re-architect session 1's ending around **"tomorrow's five minutes"**: a learned-today recap, a named tomorrow, a chosen reminder slot, and a day-2 card that opens the promised lesson. Email-first to de-risk deliverability; push and streaks as the next quarter's compounding layer.

The proof

Pre-registered 50/50 experiment on new-user D1 live by week 4. Guardrails protect session 1. We'll know within days of launch whether the hand-off moves the number — and the funnel events tell us exactly where to tune next.

A refined-clickable prototype of this loop ships alongside this deck.